
Appendix: Self-Supervised Contrastive Pre-Training for Multivariate Temporal Point Processes

Anonymous Author(s)

Affiliation
Address
email

1 1 Synthetic Generators

2 We generated datasets from Hawkes process and Proximal Graphical Event Model (PGEM). We
3 describe the parameters used in our experiments to generate event datasets.

4 **Hawkes-Exp.** We use a standard library to simulate 10-dimensional event sequences from
5 the Hawkes dynamics with exponential decay¹. The parameters which control the generation are
6 baseline rate, decay coefficient, adjacency (infectivity matrix) and end time. The four parameters
7 for models A, B, C, D, E and F in our study are different from each other, and thus represent a
8 variety of generating mechanisms which are implied by differing baseline rates, decaying patterns,
9 triggering behaviors and sequence time horizons. The actual values for the 6 models used in our
10 experiments are listed in the following for the sake of reproducibility.

11
12 **A.** baseline = [0.1097627 , 0.14303787, 0.12055268, 0.10897664, 0.08473096, 0.12917882,
13 0.08751744, 0.1783546 , 0.19273255, 0.0766883], decay = 2.5, infectivity = [[0.15037453,
14 0.10045448, 0.10789028, 0.17580114, 0.01349208, 0.01654871, 0.00384014, 0.1581418 ,
15 0.14779747, 0.16524382], [0.1858717 , 0.1517864 , 0.08765006, 0.14824807, 0.02246419,
16 0.12154197, 0.02722749, 0.17942359, 0.0991161 , 0.07875789], [0.05024778, 0.14705235,
17 0.0866379 , 0.10796424, 0.0035688 , 0.11730922, 0.11625704, 0.11717598, 0.17924869,
18 0.12950002], [0.06828233, 0.08300669, 0.13250303, 0.01143879, 0.12664085, 0.12737611,
19 0.03995854, 0.02448733, 0.05991018, 0.0690806], [0.10829905, 0.0833048 , 0.18772458,
20 0.01938165, 0.03967254, 0.03063796, 0.12404668, 0.04810838, 0.0885677 , 0.04642443],
21 [0.03019353, 0.02096386, 0.1246585 , 0.02624547, 0.03733743, 0.07003299, 0.15593352,
22 0.01844271, 0.1591532 , 0.01825224], [0.18546165, 0.08901222, 0.18551894, 0.11487999,
23 0.14041038, 0.00744305, 0.05371431, 0.02282927, 0.05624673, 0.02255028], [0.06039543,
24 0.07868212, 0.01218371, 0.13152315, 0.10761619, 0.05040616, 0.09938195, 0.01784238,
25 0.10939112, 0.1765038], [0.06050668, 0.1267631 , 0.02503273, 0.13605401, 0.0549677 ,
26 0.03479404, 0.11139803, 0.00381908, 0.15744288, 0.00089182], [0.12873957, 0.05128336,
27 0.13963744, 0.18275114, 0.04724637, 0.10943116, 0.11244817, 0.10868939, 0.04237051,
28 0.18095826]] and end-time = 10.

29 **B.** Baseline = [8.34044009e-02, 1.44064899e-01, 2.28749635e-05, 6.04665145e-02, 2.93511782e-
30 02, 1.84677190e-02, 3.72520423e-02, 6.91121454e-02, 7.93534948e-02, 1.07763347e-01], decay
31 = 2.5, adjacency = [[0. , 0. , 0.03514877, 0.15096311, 0. , 0.11526461, 0.07174169, 0.09604815,
32 0.02413487, 0.], [0. , 0. , 0. , 0.15066599, 0.15379789, 0.01462053, 0.00671417, 0. ,
33 0.], [0.01690747, 0.07239546, 0.16467728, 0. , 0.11894528, 0. , 0.11802102, 0.14348615,
34 0.00314406, 0.12896239], [0.17000181, 0. , 0. , 0. , 0. , 0. , 0.0504772 , 0.04947341,
35 0.], [0. , 0.11670321, 0.03638242, 0. , 0. , 0.00917392, 0. , 0.0252251 , 0.10131151,
36 0.1203002], [0. , 0.07118317, 0.11937903, 0.07120436, 0. , 0. , 0. , 0.08851807, 0.16239168,

¹<https://x-datainitiative.github.io/tick/>

37 0.10083865], [0. , 0.02363421, 0.02394394, 0.1388041 , 0.06836732, 0.02842716, 0.15945428,
 38 0. , 0. , 0.12481123], [0.15185513, 0. , 0.1290996 , 0. , 0. , 0.15401787, 0. , 0.16587219,
 39 0.11405672, 0.10687992], [0. , 0. , 0.07734744, 0.09943488, 0.07016556, 0.04074891, 0. ,
 40 0. , 0.00049346, 0.10609756], [0.05615574, 0.09061013, 0.15230831, 0.06142066, 0.15619243,
 41 0.10716606, 0.00271994, 0.15978585, 0.11877677, 0.17145652]] and end-time = 40.
 42 **C.** baseline = [0.08719898, 0.00518525, 0.1099325 , 0.08706448, 0.08407356, 0.06606696,
 43 0.04092973, 0.12385419, 0.05993093, 0.05336546], decay = 2.5, adjacency = [[0. , 0.09623737, 0.
 44 , 0. , 0. , 0.14283231, 0. , 0. , 0. , 0.], [0. , 0. , 0. , 0. , 0. , 0. , 0.], [0. , 0. , 0. , 0. ,
 45 0. , 0.14434229, 0.10548788, 0. , 0.], [0. , 0. , 0.16178088, 0. , 0. , 0. , 0. , 0.], [0. , 0. , 0.
 46 , 0. , 0. , 0. , 0. , 0.], [0.14554646, 0. , 0. , 0. , 0.09531432, 0. , 0. , 0. , 0.], [0. , 0.
 47 , 0. , 0. , 0. , 0. , 0. , 0.04899491], [0. , 0. , 0. , 0.07048057, 0. , 0.07546073, 0. , 0. , 0.
 48 , 0.], [0.05697369, 0. , 0. , 0. , 0. , 0.11709833, 0. , 0.03100542, 0.], [0. , 0. , 0. , 0. , 0.
 49 0.04016475, 0. , 0.10768499, 0.06297179]] and end-time = 80.
 50 **D.** Baseline = [0.11015958, 0.14162956, 0.05818095, 0.10216552, 0.17858939, 0.17925862,
 51 0.02511706, 0.04144858, 0.01029344, 0.08816197], decay = [[8., 9., 2., 7., 3., 3., 2., 4., 6., 9.],
 52 [2., 9., 8., 9., 2., 1., 6., 5., 2., 6.], [5., 8., 7., 1., 3., 5., 6., 9., 9.], [8., 6., 2., 2., 6., 6., 8., 5., 4.],
 53 [1., 1., 1., 3., 3., 8., 1., 6., 1.], [2., 5., 2., 3., 3., 5., 9., 1., 7., 1.], [5., 2., 6., 2., 9., 9., 8., 1., 1., 2.],
 54 [8., 9., 8., 5., 1., 1., 5., 4., 1., 9.], [3., 8., 3., 2., 4., 3., 5., 2., 3., 3.], [8., 4., 5., 2., 7., 8., 2., 1., 1.,
 55 6.]], adjacency = [[0. , 0.14343731, 0. , 0.11247478, 0. , 0. , 0. , 0.09416725, 0. , 0.], [0. , 0. ,
 56 0. , 0. , 0. , 0. , 0.15805746, 0. , 0.08262718], [0. , 0. , 0.03018515, 0. , 0. , 0. , 0. , 0. ,
 57 0.], [0. , 0.11090719, 0.08158544, 0. , 0.11109462, 0. , 0. , 0. , 0. , 0.], [0. , 0. , 0.14264684,
 58 0. , 0. , 0.11786607, 0. , 0. , 0.01101593, 0.], [0.04954495, 0. , 0. , 0. , 0. , 0.12385743, 0. , 0.
 59 , 0.02375575, 0.05345351], [0. , 0.14941748, 0.02618691, 0. , 0.13608937, 0. , 0.06263167, 0. ,
 60 0.04097688, 0.14101171], [0. , 0.11902986, 0. , 0.04889382, 0. , 0. , 0.01569298, 0.03678315, 0. ,
 61 0.], [0.05359555, 0. , 0. , 0.09188512, 0. , 0.14255311, 0. , 0. , 0. , 0.], [0.12792415, 0.05843994,
 62 0.16156482, 0.11931973, 0. , 0.00774966, 0.00947755, 0. , 0. , 0.]], and end-time = 50.
 63 **E.** Baseline = [0.2, 0.2, 0.2, 0.2, 0.2, 0.2, 0.2, 0.2, 0.2, 0.2], decay = [[9., 4., 9., 9., 1., 6., 4., 6., 8.,
 64 7.], [1., 5., 8., 9., 2., 7., 3., 3., 2., 4.], [6., 9., 2., 9., 8., 9., 2., 1., 6., 5.], [2., 6., 5., 8., 7., 1., 1., 3., 5.,
 65 6.], [19., 9., 8., 6., 2., 2., 6., 6., 8.], [5., 4., 1., 1., 1., 1., 3., 3., 8., 1.], [6., 1., 2., 5., 2., 3., 3., 5., 9.,
 66 1.], [7., 1., 5., 2., 6., 2., 9., 9., 8., 1.], [1., 2., 8., 9., 8., 5., 1., 1., 5., 4.], [1., 9., 3., 8., 3., 2., 4., 3.,
 67 5., 2.]], adjacency= [[0.02186539, 0.09356695, 0. , 0.16101355, 0.11527002, 0.09149395, 0. , 0. ,
 68 0.15672219, 0.], [0. , 0. , 0.14241135, 0. , 0.11167029, 0. , 0. , 0. , 0.0934937 , 0.], [0. , 0. , 0.
 69 , 0. , 0. , 0. , 0. , 0.15692693, 0.], [0.08203618, 0. , 0. , 0.02996925, 0. , 0. , 0. , 0. , 0. , 0.
 70], [0. , 0. , 0.11011391, 0.08100189, 0. , 0.11029999, 0. , 0. , 0. , 0.], [0. , 0. , 0. , 0.14162654,
 71 0. , 0. , 0.11702301, 0. , 0. , 0.01093714], [0. , 0.04919057, 0. , 0. , 0. , 0.12297152, 0. , 0.
 72 , 0.02358583], [0.05307117, 0. , 0.14834875, 0.02599961, 0. , 0.13511597, 0. , 0.06218368, 0. ,
 73 0.04068378], [0.1400031 , 0. , 0.11817848, 0. , 0.0485441 , 0. , 0. , 0.01558073, 0.03652006, 0.],
 74 [0. , 0.05321219, 0. , 0. , 0.0912279 , 0. , 0.14153347, 0. , 0. , 0.]], and end-time = 50.
 75 **F.** Baseline = [0.21736198, 0.11134775, 0.16980704, 0.33791045, 0.00188754, 0.04862765,
 76 0.26829963, 0.3303411 , 0.05468264, 0.23003733], decay = [[5., 1., 7., 3., 5., 2., 6., 4., 5., 5.],
 77 [4., 8., 2., 2., 8., 8., 1., 3., 4., 3.], [6., 9., 2., 1., 8., 7., 3., 1., 9., 3.], [6., 2., 9., 2., 6., 5., 3., 9.,
 78 4., 6.], [1., 4., 7., 4., 5., 8., 7., 4., 1., 5.], [5., 6., 8., 7., 7., 3., 5., 3., 8., 2.], [7., 7., 1., 8., 3., 4.,
 79 6., 5., 3., 5.], [4., 8., 1., 1., 6., 7., 6., 7., 5.], [8., 4., 3., 4., 9., 8., 2., 6., 4., 1.], [7., 3., 4., 5.,
 80 9., 9., 6., 3., 8., 6.]], adjacency = [[0.15693854, 0.04896059, 0.0400508 , 0. , 0. , 0.13021228,
 81 0. , 0.10699903, 0. , 0.15329807], [0.0784283 , 0. , 0. , 0. , 0. , 0. , 0.00310706, 0.0090892 ,
 82 0.07758874, 0.], [0.01672489, 0. , 0. , 0.07851303, 0. , 0. , 0.12848331, 0.08859293, 0. , 0.],
 83 [0.09984995, 0. , 0. , 0.10541925, 0. , 0.08032527, 0. , 0. , 0. , 0.], [0.14642469, 0.06629365,
 84 0. , 0. , 0. , 0. , 0.11891738, 0.04166225, 0.09808829, 0.17638655], [0.00976324, 0.1100343 ,
 85 0.02003261, 0. , 0. , 0. , 0.05993539, 0.09739541, 0. , 0. , 0.], [0. , 0. , 0. , 0. , 0.04672133, 0.16916
 86 , 0. , 0.17341419, 0.12078975, 0.14441602, 0.], [0. , 0.17305542, 0.06927975, 0. , 0.03408974,
 87 0. , 0.08457162, 0.03787486, 0.10863292, 0.], [0.07186225, 0.05760593, 0. , 0. , 0.08042031,
 88 0.04403479, 0.1033595 , 0.17046747, 0. , 0.05083523], [0. , 0.10029222, 0. , 0.1022067 , 0. , 0. ,
 89 0.0588527 , 0. , 0.03530513, 0.]], and end-time = 40.

90 **PGEM.** We implement the PGEM generator by previous work [1] and generate 5-dimensional
 91 event datasets governed by the PGEM dynamics, which is signified by its piecewise conditional
 92 intensity function for each event type at a given time instance of a parental configuration observed

93 at a particular window. The parameters are the conditional intensity (lambdas) for each event type
 94 given parental states, parental configuration (parents), windows for each parental state (windows)
 95 and end time which is 100 across all models. The parameters for models A, B, C, D and E are
 96 listed in the following. We point out to the readers that the 5 models share different graphs, which
 97 is implied by the parents for each event type (node).

98 **A** has the following parameters: 'parents': 'A': [], 'B': [], 'C': ['B'], 'D': ['A', 'B'], 'E': ['C'],
 99 'windows': 'A': [], 'B': [], 'C': [15], 'D': [15, 30], 'E': [15], 'lambdas': 'A': () 0: 0.2, 'B': () 0: 0.05,
 100 'C': (0,): 0.2, (1,): 0.3, 'D': (0, 0): 0.1, (0, 1): 0.05, (1, 0): 0.3, (1, 1): 0.2, 'E': (0,): 0.1, (1,): 0.3

101 **B** has the follow parameters: 'parents': 'A': ['B'], 'B': ['B'], 'C': ['B'], 'D': ['A'], 'E': ['C'],
 102 'windows': 'A': [15], 'B': [30], 'C': [15], 'D': [30], 'E': [30], 'lambdas': 'A': (0,): 0.3, (1,): 0.2,
 103 'B': (0,): 0.2, (1,): 0.4, 'C': (0,): 0.4, (1,): 0.1, 'D': (0,): 0.05, (1,): 0.2, 'E': (0,): 0.1, (1,): 0.3.

104 **C** has the follow parameters: 'parents': 'A': ['B', 'D'], 'B': [], 'C': ['B', 'E'], 'D': ['B'], 'E': ['B'],
 105 'windows': 'A': [15, 30], 'B': [], 'C': [15, 30], 'D': [30], 'E': [30], 'lambdas': 'A': (0, 0): 0.1, (0,
 106 1): 0.05, (1, 0): 0.3, (1, 1): 0.2, 'B': () 0: 0.2, 'C': (0, 0): 0.2, (0, 1): 0.05, (1, 0): 0.4, (1, 1): 0.3, 'D':
 107 (0,): 0.1, (1,): 0.2, 'E': (0,): 0.1, (1,): 0.4.

108 **D** has the follow parameters: 'parents': 'A': ['B'], 'B': ['C'], 'C': ['A'], 'D': ['A', 'B'], 'E': ['B',
 109 'C'], 'windows': 'A': [15], 'B': [30], 'C': [15], 'D': [15, 30], 'E': [30, 15], 'lambdas': 'A': (0,):
 110 0.05, (1,): 0.2, 'B': (0,): 0.1, (1,): 0.3, 'C': (0,): 0.4, (1,): 0.2, 'D': (0, 0): 0.1, (0, 1): 0.3, (1, 0):
 111 0.05, (1, 1): 0.2, 'E': (0, 0): 0.1, (0, 1): 0.02, (1, 0): 0.4, (1, 1): 0.1.

112 **E** has the follow parameters: 'parents': 'A': ['A'], 'B': ['A', 'C'], 'C': ['C'], 'D': ['A', 'E'], 'E': ['C',
 113 'D'], 'windows': 'A': [15], 'B': [30, 30], 'C': [15], 'D': [15, 30], 'E': [15, 30], 'lambdas': 'A':
 114 (0,): 0.1, (1,): 0.3, 'B': (0, 0): 0.01, (0, 1): 0.05, (1, 0): 0.1, (1, 1): 0.5, 'C': (0,): 0.2, (1,): 0.4, 'D':
 115 (0, 0): 0.05, (0, 1): 0.02, (1, 0): 0.2, (1, 1): 0.1, 'E': (0, 0): 0.1, (0, 1): 0.01, (1, 0): 0.3, (1, 1): 0.1.

116 2 Real Datasets

117 **Cosmetics and Electronics** contain user-level online transactions in an electronics and cosmetics
 118 store respectively. While the original cosmetics dataset from Kaggle contain multiple months of
 119 transactions, we used the one from Dec, 2019. Similarly, we also used electronics dataset from Dec,
 120 2019. Both share the same four types of events: 'view', 'cart', 'remove-from-cart' and 'purchase'. The
 121 two datasets involve transaction events in seconds. To optimize computation for transformer models,
 122 we filtered out sequences longer than 300 events and shorter than 30 and scaled the timestamps
 123 into [0,1] to avoid numerical issues. We test the algorithms on Electronics after pre-training on
 124 Cosmetics.

125 **Defi-Mainnet and Defi-Polygon.** Defi Mainnet is built on the more widely used ethereum
 126 blockchain. Polygon is a scalable sidechain of Ethereum that allows for much faster and low fee
 127 transactions than the original Ethereum Blockchain. The difference in fee structure produces quite
 128 different dynamics in the two different AAVE lending protocols. DeFi-Polygon has many more
 129 users and transactions per user, but much less total value locked than Defi Mainnet. The polygon
 130 users are much more likely to engage in risky but potentially profitable "Yield Farming" transac-
 131 tions. Foundation methods would be very useful in modeling the many other lending protocols in
 132 AAVE or other lending platforms which are new or less popular and thus have less transactions.
 133 The 6 types of actions user performs are : 'borrow', 'collateral', 'deposit', 'liquidation', 'redeem',
 134 'repay'. The origin timestamps are mined in Unix Timestamp. We filtered out sequences longer than
 135 300 events and shorter than 30 and scaled the timestamps into [0,1] to avoid numerical issues. We
 136 test the algorithms on Mainnet after pre-training on Polygon.

137 **ACLED-India and ACLED-Bangladesh** contains sequences where each sequence involves an ac-
 138 tor involving some armed conflict (i.e. riots and protests) in the respective country. The former
 139 involves events happening from 2016 to 2022, and the latter 2010-2021. The unit of each timestamp
 140 is 'days'. There are 6 types of events in the ACLED-India which are: 'Battles', 'Explosions/Remote
 141 violence', 'Riots', 'Violence against civilians', 'Protests', and 'Strategic developments'. The 4 over-
 142 lapping types in ACLED-Bangladesh are 'Battles', 'Explosions/Remote violence', 'Riots', and 'Vi-
 143 olence against civilians'. We filtered out sequences longer than 300 events and shorter than 2 and
 144 scaled the timestamps into [0,1] to avoid numerical issues. We test the algorithms on Bangladesh
 145 after pre-training on India.

146 A descriptive summary of the above datasets are shown in Table 1.

Table 1: Properties of 6 real datasets.

Dataset	# classes	# seqs.	Avg. length	# events	Data Type
Defi-Mainnet	6	20539	32	654844	Financial
Defi-Polygon	6	33597	85	2856453	Financial
Electronics	4	9993	20	195726	E-Commerce
Cosmetics	4	19301	39	752109	E-Commerce
ACLED-India	6	111	17	1934	Political
ACLED-Bangladesh	4	97	17	1697	Political

147 3 Model Implementation and (Pre)Training

148 **Pretraining.** Our pretraining model is based on the code adaptation from [2], which can be found
 149 in the Supplementary Material. The entire procedure is detailed in Algorithm 1 where a combined
 150 loss $\mathcal{L}_{event} = \mathcal{L}_{predict} + \lambda\mathcal{L}_{contrast}$ is minimized for each batch where $\mathcal{L}_{predict}, \mathcal{L}_{contrast}$ are
 151 from equation 3 and 4 respectively. We first randomly sample void event instances in-between
 152 every consecutive real events and then insert such void events. Then for each augment sequence,
 153 we randomly mask 15% of the events and split into train and dev subsets of sequences. We train
 154 our model using stochastic gradient descent and employ the Adam optimizer for optimization. The
 155 default transformer architecture used for pretraining includes the following specifications: the multi-
 156 headed self-attention module consists of 4 blocks, the dimension of the value vector after attention
 157 is applied is 512, the number of attention heads is 4, the hidden layer of the feed-forward neural
 158 network has a dimension of 1024, the value vector has a dimension of 512, the key vector has a
 159 dimension of 512, and the dropout rate is set at 0.1. Our training process spans 100 epochs with a
 160 learning rate of 0.0001. In all experiments, the hyper-parameters γ and λ are set to 1, and 0.01. The
 161 temporal parameter, τ , is fixed at 1. It is worth pointing out that equation 4 can be simplified to a
 162 more concise form by absorbing τ into λ in practice. The outputs of the transformer model, denoted
 163 as \mathbf{H}_i 's, are linearly mapped to numerical values and probabilities in the regression and classification
 164 losses presented in equation 3.

165 **Fine-tuning.** For the fine-tuning process, we employ a feed-forward network consisting of three
 166 hidden layers, each with a dimension of 512. The representations of the training and development
 167 subsets are acquired from the pre-trained model and are held constant throughout the fine-tuning
 168 stage. We utilize the Adam optimizer with a learning rate chosen from the range of 0.001, 0.002,
 169 along with a batch size of 32, for a total of 100 epochs. To prevent overfitting, we implement early
 170 stopping if there is no improvement in the loss L_{event} for 10 consecutive epochs. The learning
 171 rate is determined based on the model's performance on the development subset of the target data.
 172 Furthermore, in all experiments, the value of α in equation 5 is fixed at 0.01.

173 **Combined Position Encoding (PE) and Temporal Encoding (TE).** We present a straightforward
 174 yet illustrative example demonstrating the effectiveness of the combined encoding scheme, which

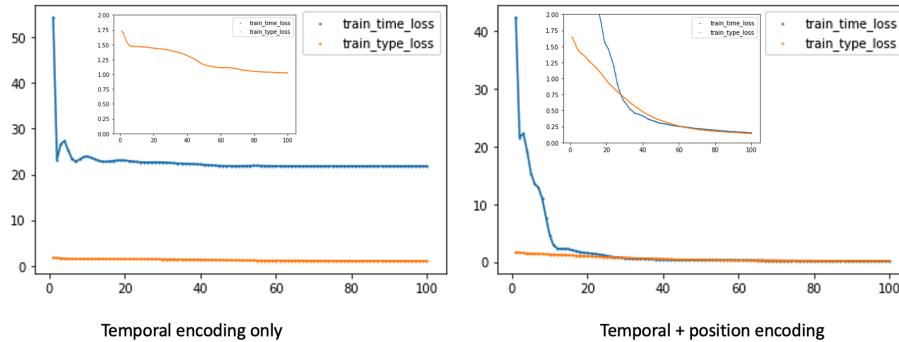


Figure 1: The effect of combined TE + PE in training.

Algorithm 1: Pretraining of Event-former

Input: Given dataset S with D sequences, each sequence S_d with length d_l , $\{(t_i, y_i)\}_{i=1}^{d_l}$, batch size b

```
 $S' = []$ 
for  $d \leftarrow 1$  to  $D$  do
| seq' = []
| for  $i \leftarrow 1$  to  $d_l - 1$  do
| |  $t'_i \sim Unif(t_i, t_{i+1})$ 
| | seq'.append(( $t'_i$ , null))
| end
| seqnew = merge( $S_d$ , seq')
| S'.append( seqnew )
end
for  $d \leftarrow 1$  to  $D$  do
| for  $i \leftarrow 1$  to  $d'_l - 1$  do
| | mask  $\sim Bern(0.15)$ 
| | if mask == 1 then
| | | (0, null)  $\leftarrow (t'_i, y'_i)$ 
| | end
| end
end
split( $S'_{m, tr}$ ) :=  $S'_{m, tr} = \{S'_{m, k}\}_{k=1}^K, S'_{m, dev}$ 
for epoch  $\leftarrow 1$  to  $N$  do
| for iteration  $\leftarrow 1$  to  $\lceil \frac{K}{b} \rceil$  do
| |  $B' \sim S'_{m, tr}$ ; Sample a batch of sequences
| | compute  $\mathcal{L}_{event}(B')$  (Eq. 3)
| | back-propagate with gradient  $\nabla_{\theta, \phi} \mathcal{L}_{event}(B')$  ( $\theta$ : Transformer model parameters,  $\phi$ : Weights and bias for regression and classification.)
| | update parameters of network  $\theta, \phi$ 
| end
| evaluate  $L_{event}(S'_{m, dev})$ , stop training if not improving in 5 epochs
end
Return: Optimal parameters  $\phi^*, \theta^*$ 
```

175 incorporates both position encoding and temporal encoding. In our experiment, we trained the model
176 with only two samples, and the results clearly demonstrate improved overall training compared
177 to using temporal encoding alone. Figure 1 visually depicts the compromised optimization when
178 employing only temporal encoding.

179 **4 Baselines Models and Implementation**

180 In our study, we provide a brief overview of the implementation of the baseline models.
181 T-RMTPP and T-ERPP: The original Recurrent Marked Temporal Point Process (RMTPP) and
182 Event Recurrent Point Process (ERPP) models are based on recurrent neural networks (RNNs) or
183 their variants. Our implementation builds upon a standard codebase ². However, to ensure a fair
184 comparison, we replaced the RNN components with transformers. We trained these models using
185 the recommended set of parameters in our experiments.
186 T-LNM: The original Lognormal Mixture (LNM) model, like the previous models, is RNN-based.
187 To model the density of log (inter-event) times, we utilized 64 mixtures of lognormal components
188 while keeping other parameters at their default setting ³.

²<https://github.com/woshiyya/ERPP-RMTPP>

³<https://github.com/shchur/ifl-tpp>

189 THP: The Transformer Hawkes Process (THP) models were implemented using attention-based
 190 transformers. We followed the recommended parameter settings for these models in our experi-
 191 ments.⁴.

192 **5 Proof Sketch of Theorem 1: Transformer with Combined Temporal 193 Encoding and Position Encoding**

194 We consider a general type of transformer architecture as described by Yun et al. (2019) in their
 195 work on transformers. Our proof demonstrates that a transformer with combined temporal encoding
 196 and position encoding can serve as a universal approximating function for any continuous function
 197 in the context of sequence-to-sequence tasks with compact support. This proof follows a similar
 198 structure to the proof of the transformer with position encoding (Theorem 3 in [3]).

199 Without loss of generality, let's consider a sequence with timestamps t_1, t_2, \dots, t_n . We assume that
 200 t_i is an integer value for each $i \in 1, 2, 3, \dots, n$, and that $t_i < t_j$ for all $i < j$. In case any t_i values
 201 are in decimal form, we can multiply them by a constant factor without affecting the dynamics of
 202 the event sequence, thereby transforming each given timestamp to an integer.

203 We define a d -dimensional temporal encoding for the n event epochs as follows:

$$\mathbf{T} = \begin{bmatrix} 0 & t_2 - t_1 & \dots & t_n - t_1 \\ 0 & t_2 - t_1 & \dots & t_n - t_1 \\ \vdots & \vdots & & \vdots \\ 0 & t_2 - t_1 & \dots & t_n - t_1 \end{bmatrix}$$

204 Similarly, a d -dimensional position encoding for the n event epochs is defined as:

$$\mathbf{P} = \begin{bmatrix} 0 & 1 & \dots & n \\ 0 & 1 & \dots & n \\ \vdots & \vdots & & \vdots \\ 0 & 1 & \dots & n \end{bmatrix}$$

205 The combined encoding is given by:

$$\mathbf{PE} + \mathbf{TE} = \begin{bmatrix} 0 & t_2 - t_1 + 1 & \dots & t_n - t_1 + 1 \\ 0 & t_2 - t_1 + 1 & \dots & t_n - t_1 + 1 \\ \vdots & \vdots & & \vdots \\ 0 & t_2 - t_1 + 1 & \dots & t_n - t_1 + 1 \end{bmatrix}$$

206 The strict temporal ordering of the timestamps guarantees that $t_i - t_1 + 1 < t_{i+1} - t_1 + 1$ for all
 207 $i \in \{1, 2, 3, \dots, n-1\}$. This ensures that for all rows, the coordinates are monotonically increasing,
 208 and the input values can be partitioned into cubes.

209 The remaining steps of the proof closely follow those outlined in the proof of Theorem 3 in [3]. This
 210 involves replacing n with t_n through quantization performed by feed-forward layers, contextual
 211 mapping achieved by attention layers, and function value mapping implemented by feed-forward
 212 layers.

213 **6 Clarification, Limitations And Societal Impact**

214 **Additional rationale for pre-training and fine-tuning.** In the context of dynamic event datasets,
 215 the role of labels extends beyond being mere targets for prediction, as observed in static image and
 216 text classification tasks. Here, labels in dynamic event datasets also serve as dynamic "features"
 217 in addition to their target role. Self-supervised representation learning emerges as a particularly
 218 effective approach for achieving a concise representation of the dynamic event history, utilizing
 219 event labels and timestamps as raw features. Our objective is to capitalize on a large volume of event
 220 sequences to acquire high-quality representations of event dynamics, especially for downstream

⁴<https://github.com/SimiaoZuo/Transformer-Hawkes-Process>

221 tasks with limited data availability. Consequently, learning the representation of historical context
222 and transferring a sequence of such representations that encapsulate event dynamics differs from
223 learning representations for tabular, image, or text data. A similar motivation can be observed in
224 time series transfer learning [4], where self-supervised techniques demonstrate their effectiveness in
225 acquiring compact dynamic features from raw historical values in a time series.

226 **Advantages of homogeneous transfer.** The homogeneous transfer problem, where pre-training
227 on open domains and target task domains share the same set of event types, is not an oversimplified
228 scenario. In fact, a significant portion of transfer learning research focuses on this type of transfer
229 [5]. In our study, we also consider this scenario, where the event types are shared across pre-training
230 and fine-tuning; however, the datasets themselves may differ, resulting in varying underlying dynam-
231 ics between datasets. We have presented several practical applications in our paper, supported by
232 experimental evidence and thorough exposition. In the case of alignment of event types, we utilize
233 shared embedding matrices for event types during both pre-training and fine-tuning in our homo-
234 geneous transfer setting. Therefore, the challenge of aligning event types does not arise. However,
235 it is worth noting that addressing the alignment issue becomes relevant for heterogeneous transfer,
236 which is an avenue we are currently exploring.

237 **Impact of self-supervised pretraining for TPP.** Our self-supervised pretraining approach for
238 temporal point process (TPP) models has certain considerations and potential societal implications
239 that warrant attention. Firstly, the availability and representativeness of large-scale event streams,
240 upon which our pretraining relies, may vary, posing challenges in accessing diverse real-world sce-
241 narios. Additionally, self-supervised training can be computationally demanding, necessitating sub-
242 stantial computational resources and time investments. Moreover, we acknowledge the potential
243 risks associated with biases present in some data streams, which could lead to biased model predic-
244 tions and the potential reinforcement of societal inequalities. Therefore, caution should be exercised
245 when utilizing our model, and careful data curation and preprocessing are essential to mitigate such
246 biases. Furthermore, the widespread adoption of self-supervised training for TPP models has im-
247 plications for domains like healthcare and finance. To prevent potential negative consequences and
248 societal harm, ethical considerations encompassing privacy protection, fairness, and transparency
249 must be addressed. While self-supervised training holds promising opportunities, its limitations and
250 societal impact demand meticulous consideration to ensure responsible harnessing of its potential
251 and its beneficial application in real-world contexts.

252 References

- 253 [1] Debarun Bhattacharjya, Dharmashankar Subramanian, and Tian Gao. Proximal graphical event
254 models. *Advances in Neural Information Processing Systems*, 31, 2018.
- 255 [2] Simiao Zuo, Haoming Jiang, Zichong Li, Tuo Zhao, and Hongyuan Zha. Transformer Hawkes
256 process. In *International Conference on Machine Learning*, pages 11692–11702. PMLR, 2020.
- 257 [3] Chulhee Yun, Srinadh Bhojanapalli, Ankit Singh Rawat, Sashank J Reddi, and Sanjiv Kumar.
258 Are transformers universal approximators of sequence-to-sequence functions? *arXiv preprint*
259 *arXiv:1912.10077*, 2019.
- 260 [4] George Zerveas, Srideepika Jayaraman, Dhaval Patel, Anuradha Bhamidipaty, and Carsten
261 Eickhoff. A transformer-based framework for multivariate time series representation learning.
262 In *Proceedings of the 27th ACM SIGKDD Conference on Knowledge Discovery & Data Mining*,
263 pages 2114–2124, 2021.
- 264 [5] Fuzhen Zhuang, Zhiyuan Qi, Keyu Duan, Dongbo Xi, Yongchun Zhu, Hengshu Zhu, Hui
265 Xiong, and Qing He. A comprehensive survey on transfer learning. *Proceedings of the IEEE*,
266 109(1):43–76, 2021.